Architecture Core LLM Platform

CACIB Gen AI

AWS Professional Services

15th May 2024

Agenda

1. Introduction & Context
2. Logical component of a Gen AI platform
3. Requirements for wave 1 applications go to production
4. Mapping CACIB use cases with Gen AI platform
5. Design Principles
6. Gen AI Core Platform Blueprint
7. Use cases architecture proposal

Introduction & Context

Starting April 22nd, the program aims to deploy four use cases in production by September, leveraging a scalable and innovative Generative AI Core Platform on AWS.

Four discovery workshops have been conducted to capture the specific requirements of the use cases. Additional focused workshops are planned (FinOps, DevOps, Security) to delve deeper into the remaining open areas to enhance the global architecture design, ensuring that synergies and opportunities for mutualization are identified.

Through this presentation, we are aiming at aligning all the key stakeholders around :

- The definition of the Gen AI Core Platform Lot 1 for September
- Architecture principles to be validated and the associated action plan

The global alignment of all stakeholders is crucial to ensure that the various architectural documents (GAD, DAD, RPAD) are implemented within the suggested timeframe (end of May).

Logical components of a Gen AI platform

Managed by the Landing zone

Core platform capabilities

Governance dependant capability

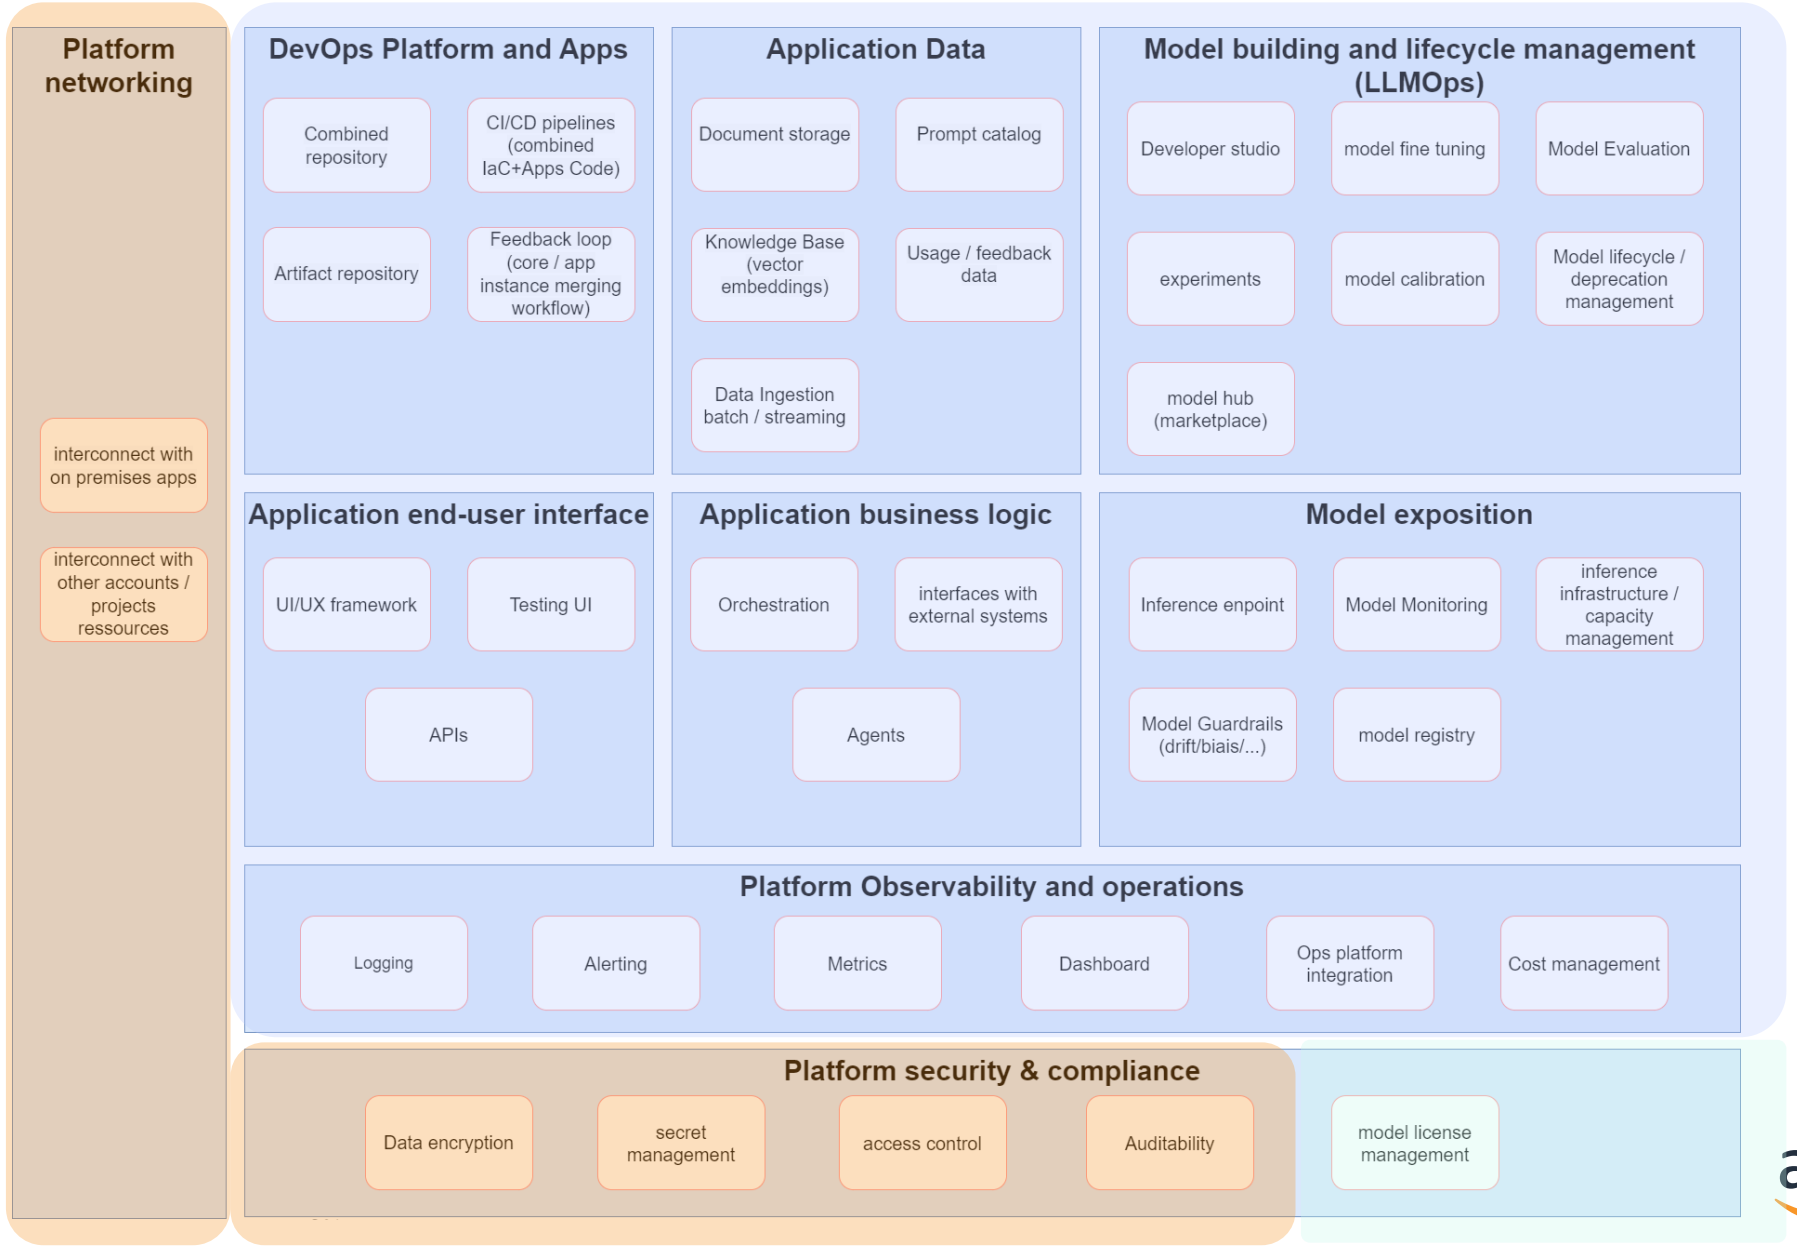

aws financial services

Requirements for wave 1 applications go to production

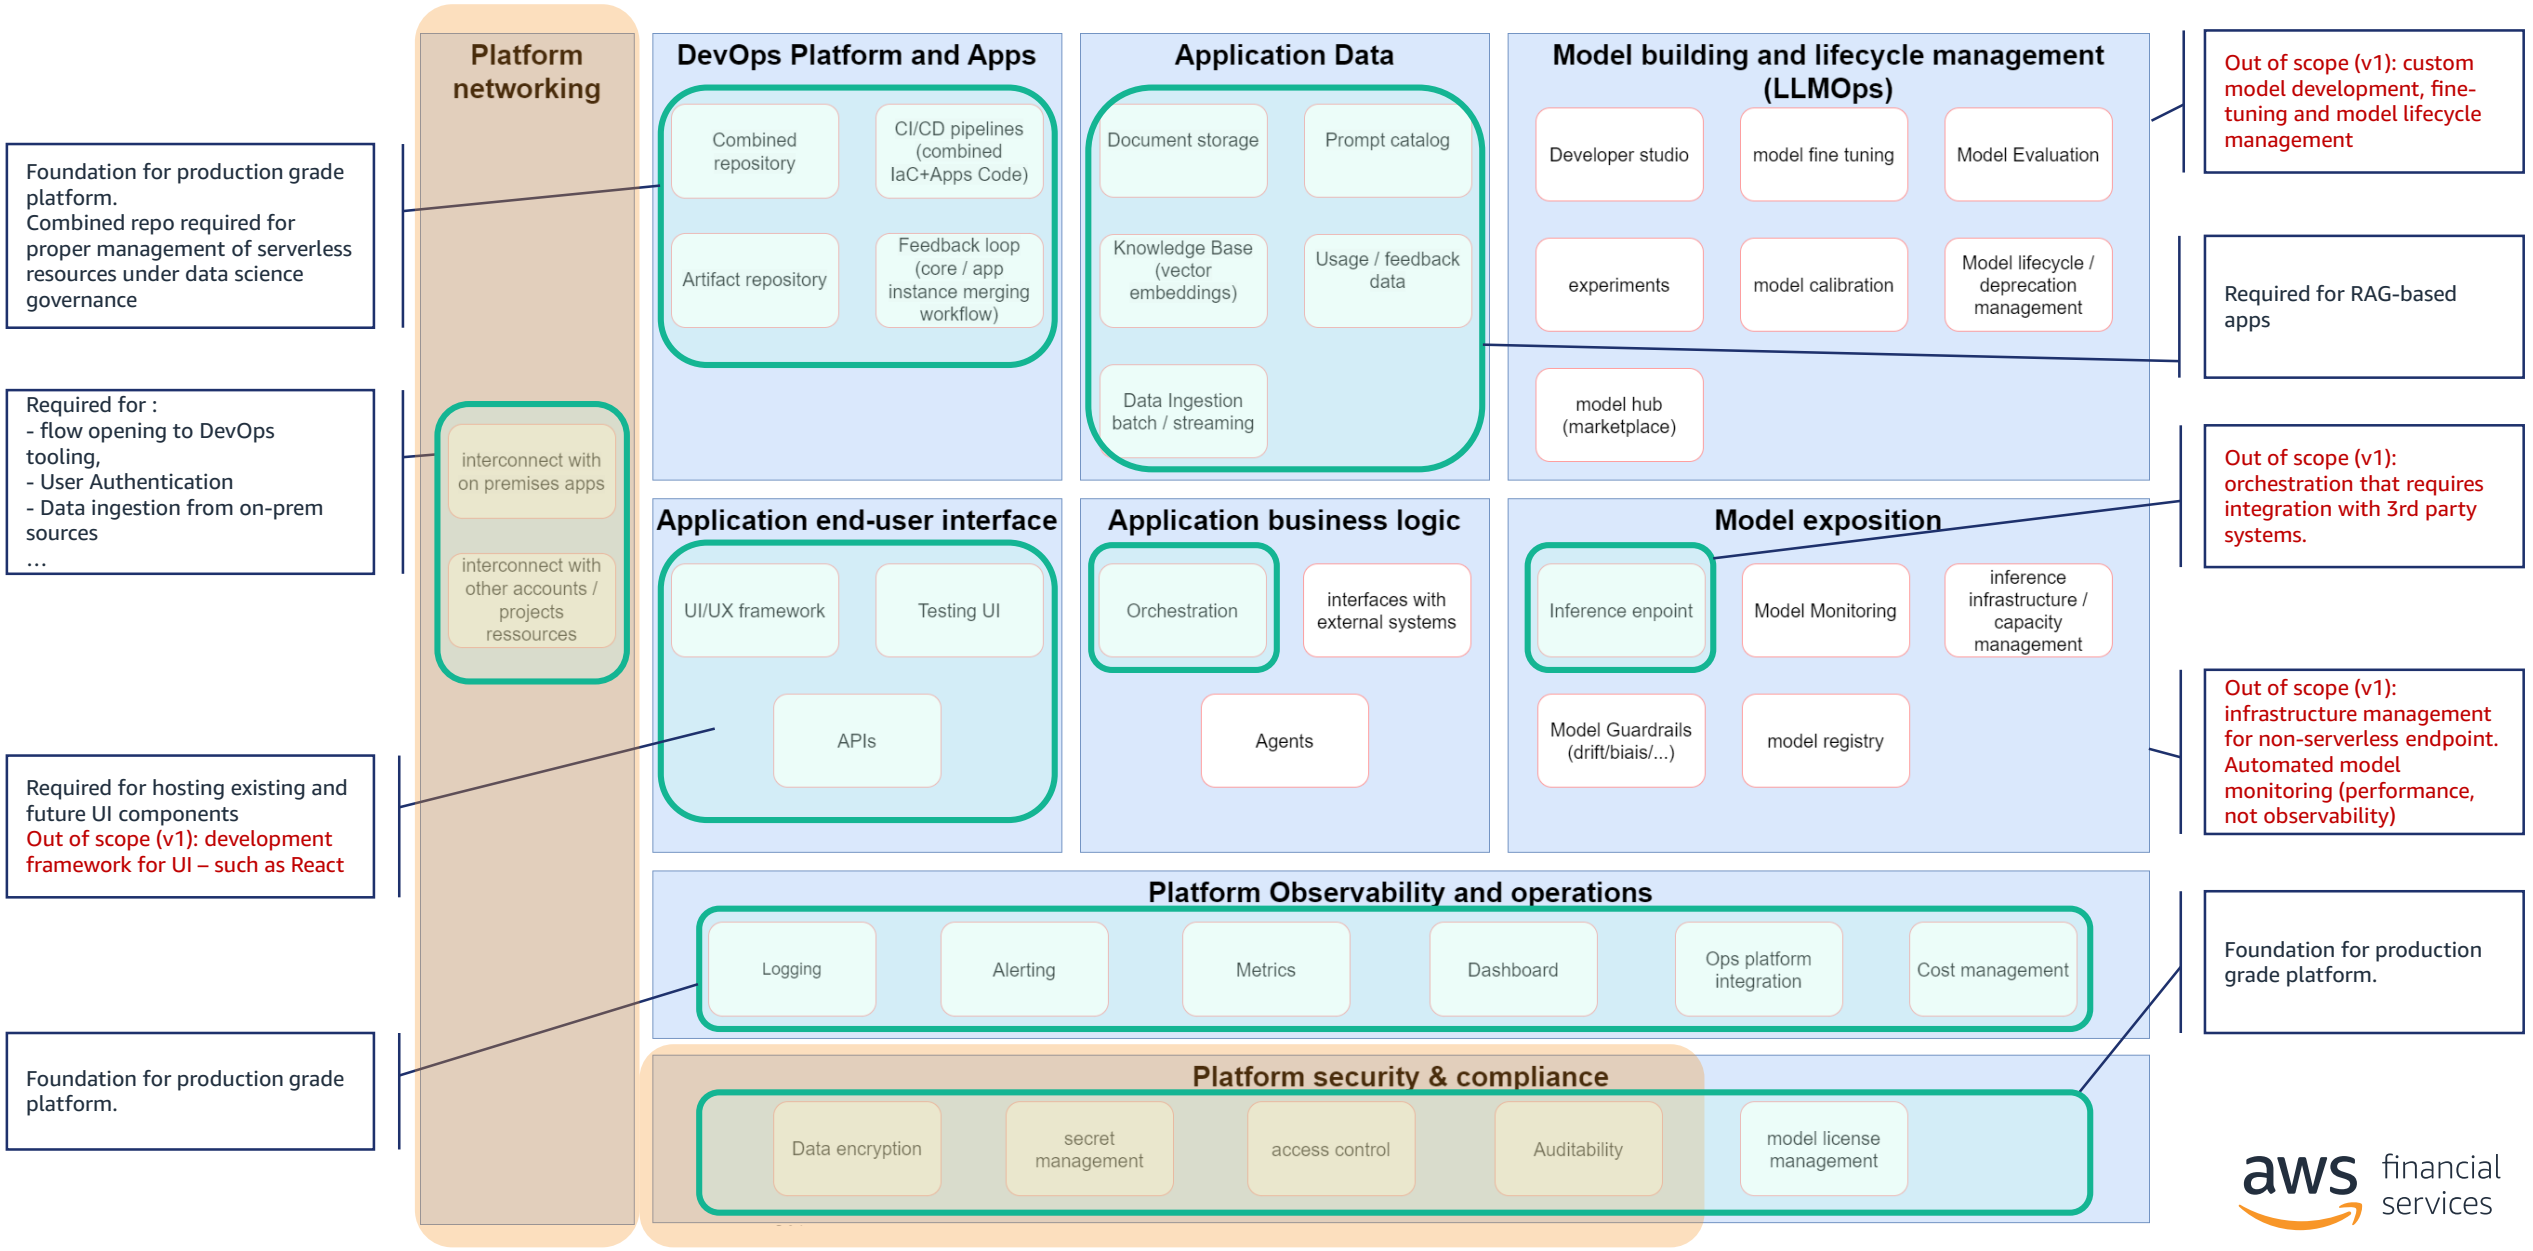

aws financial services

Mapping logical components with AWS services

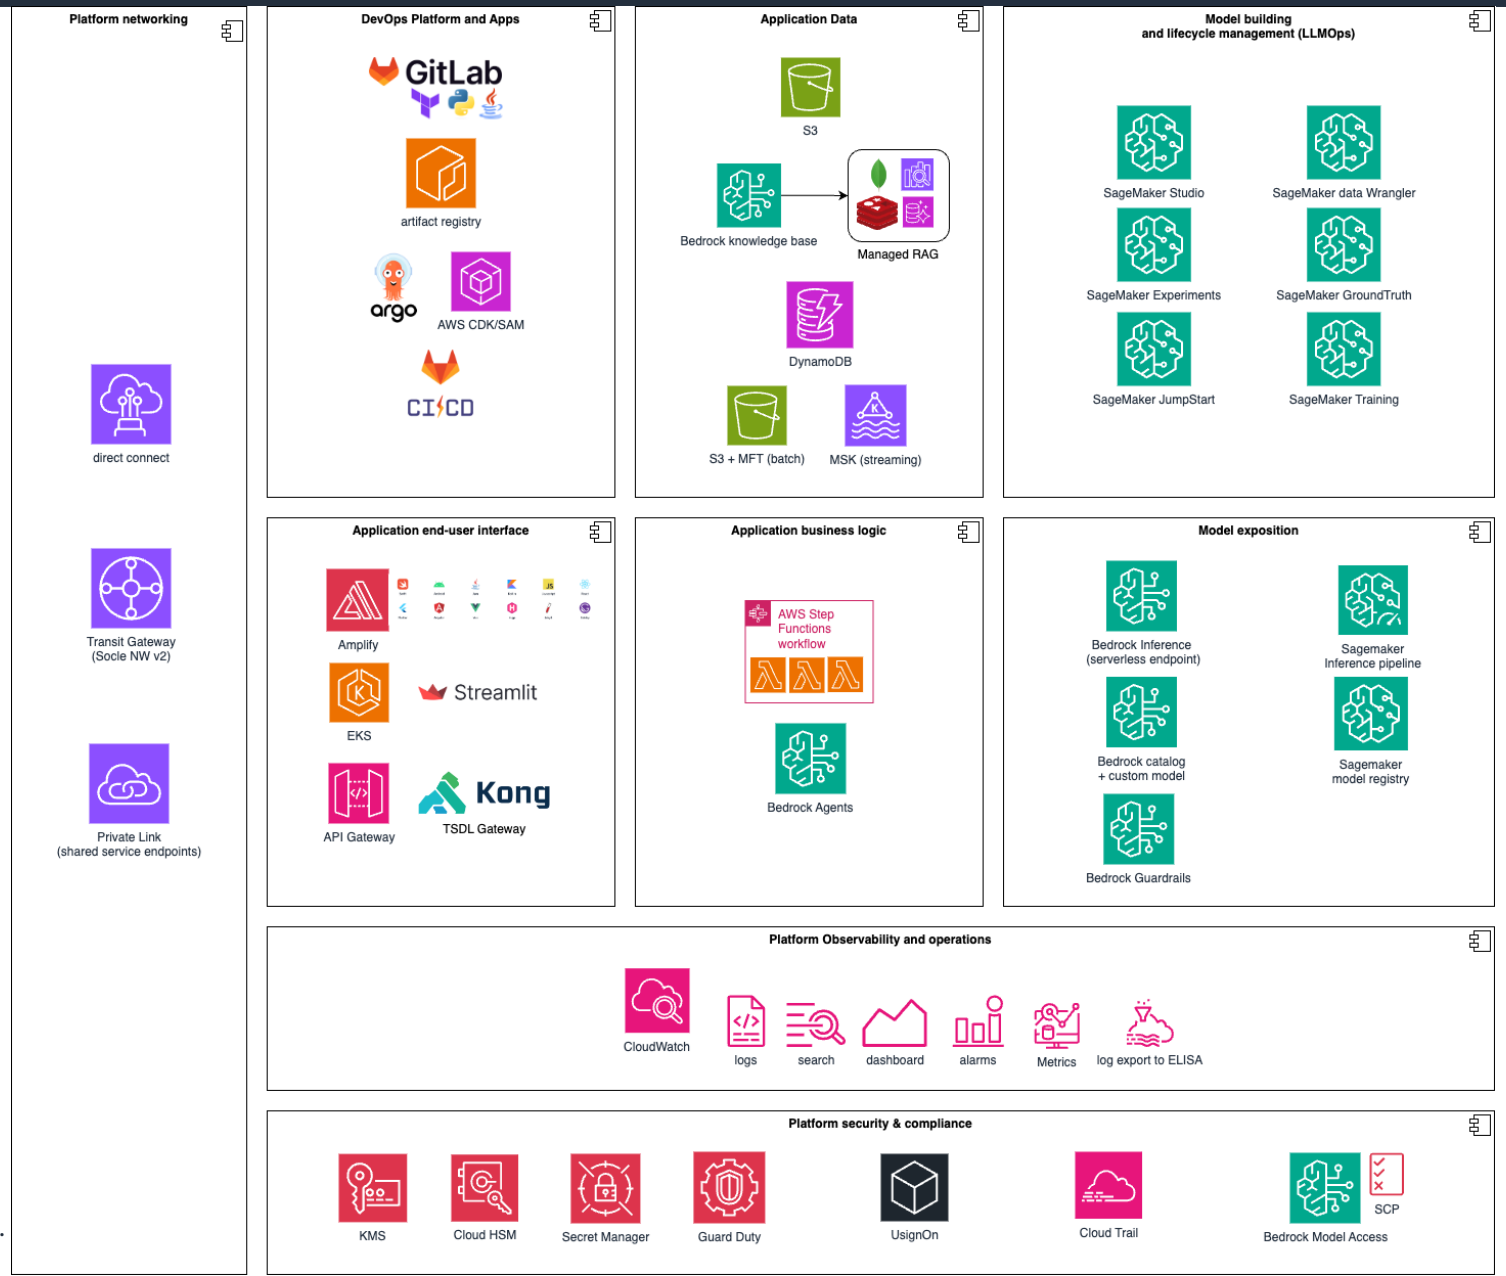

aws financial services

Design orientation principles

- Prioritize serverless services over managed or non-managed infrastructure resources to start small, iterate quickly and scale fast.
- Prefer managed over non-managed services if no restrictions stating the opposite are in place.
- Modularize functionalities and components whenever possible to improve platform evolvability and scalability.
- Define the code base and its automation processes (merge, deploy) in a single repository that acts as the source of truth for all the base components of the core platform.
- Reuse existing CA and CACIB vetted cloud architecture patterns (SSO, integration, etc) whenever possible to streamline deployment and simplify operations.
- Enforce zero trust security accross the core platform, within and across accounts.

aws financial services

What is the Core Platform ?

Based on the design principles and the architecture decisions proposed, the **Core platform is**:

- Infrastructure as code blueprint
- Re-usable modules, artifacts & patterns as you go
- Shared standards for LLM applications
- Evolutive modular & decoupled architecture
- One-to-one project and blueprint mapping (1 project = 1 deployed blueprint)
- Traceable and scalable feature addition and improvements (via PRs)
- Centralized roadmap with decentralized management and ownership
- Quality gates between promotion stages to ensure best practices

aws financial services

DevOps proposed approach

WiP - Pending for DevOps discovery workshop

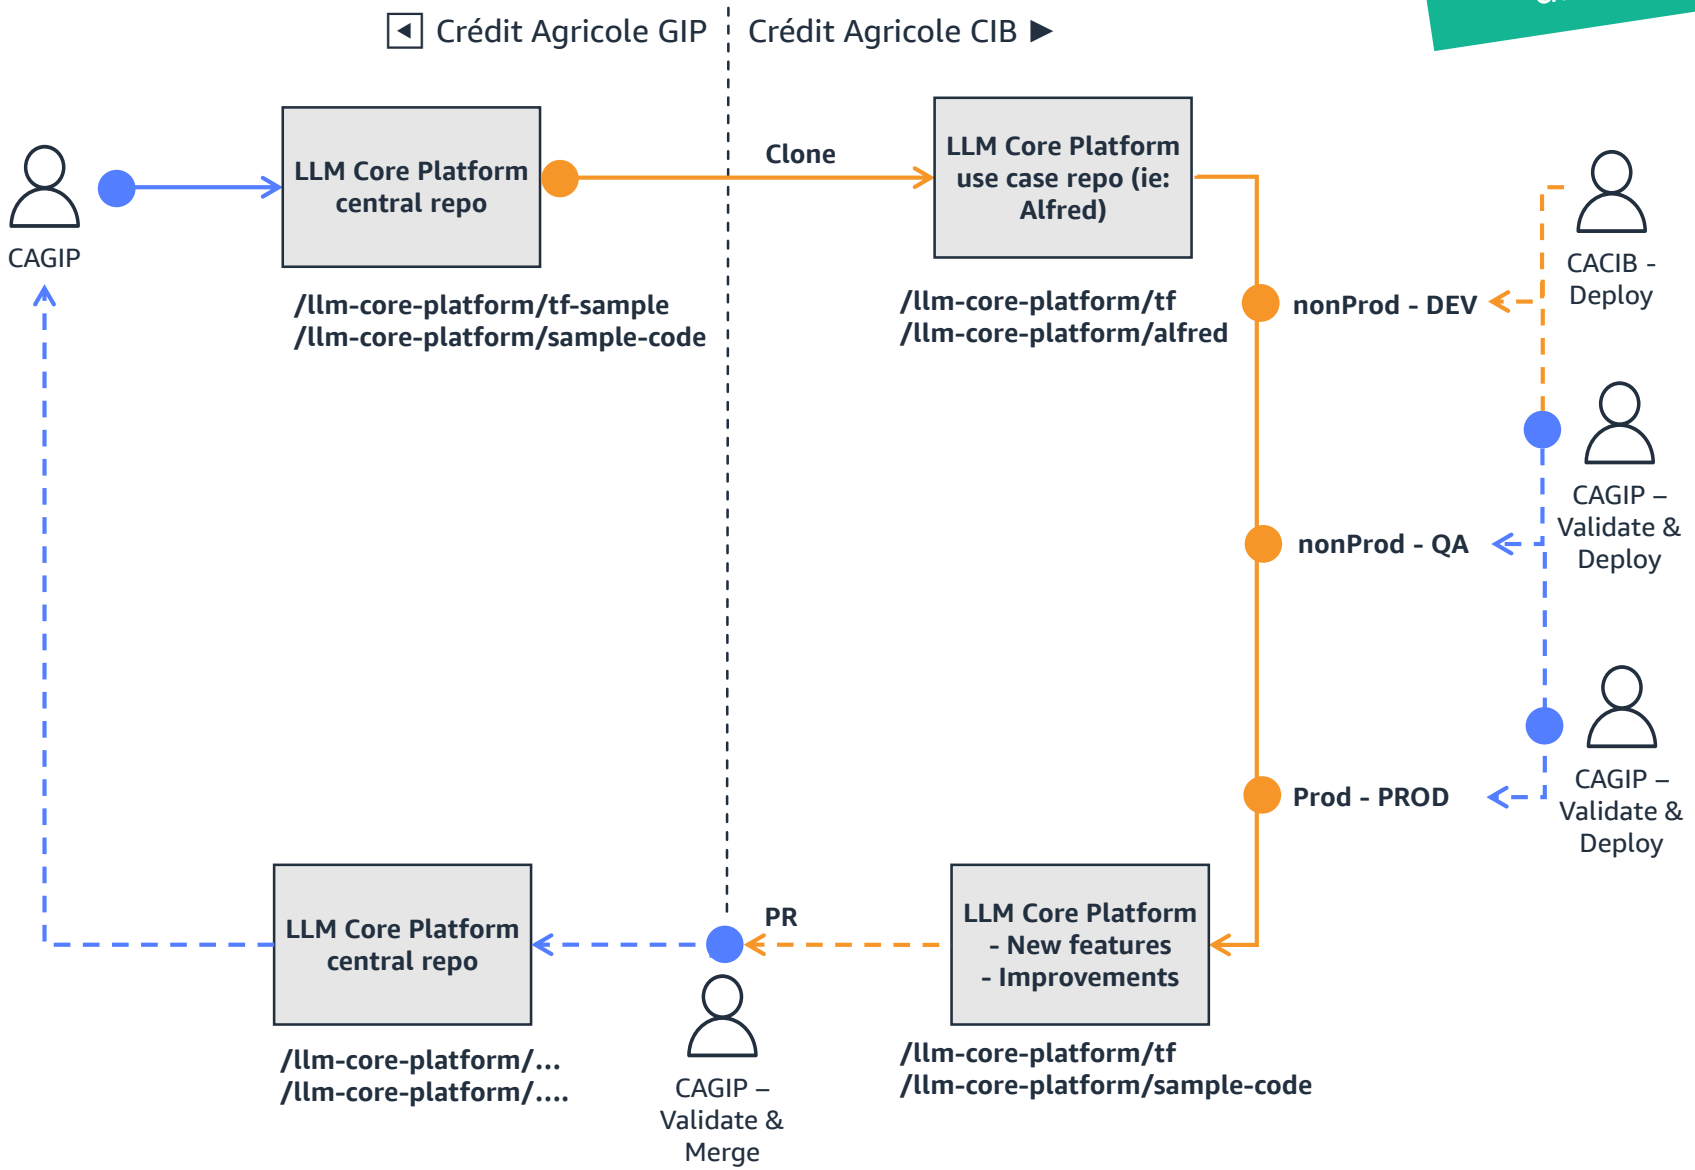

aws financial services

Decentralized platform approach

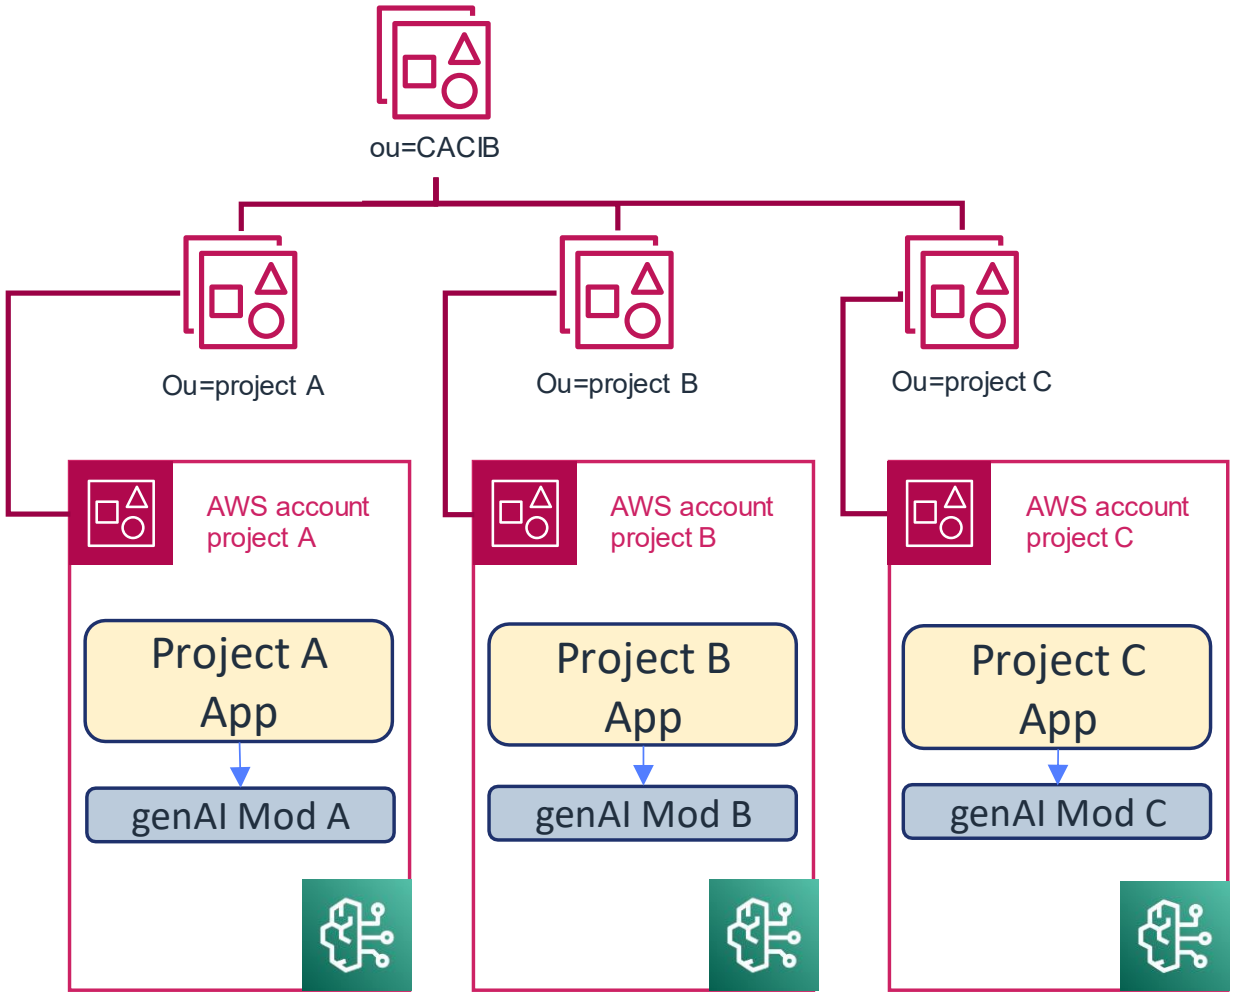

- Easier project level FinOps
- More flexibility for core model tuning
- Promote decentralized DevOps practices and empowers teams
- Facilitate reversibility through project independence
- Aligned with CACIB governance and processes for project qualification, architecture validation and security review.
- Aligned with general cloud best practices and Landing Zone usage guidelines.
- Aligned with existing CACIB projects on AWS
- No minimal deployment footprint (scale to 0)

aws financial services

Gen AI Core Platform Blueprint

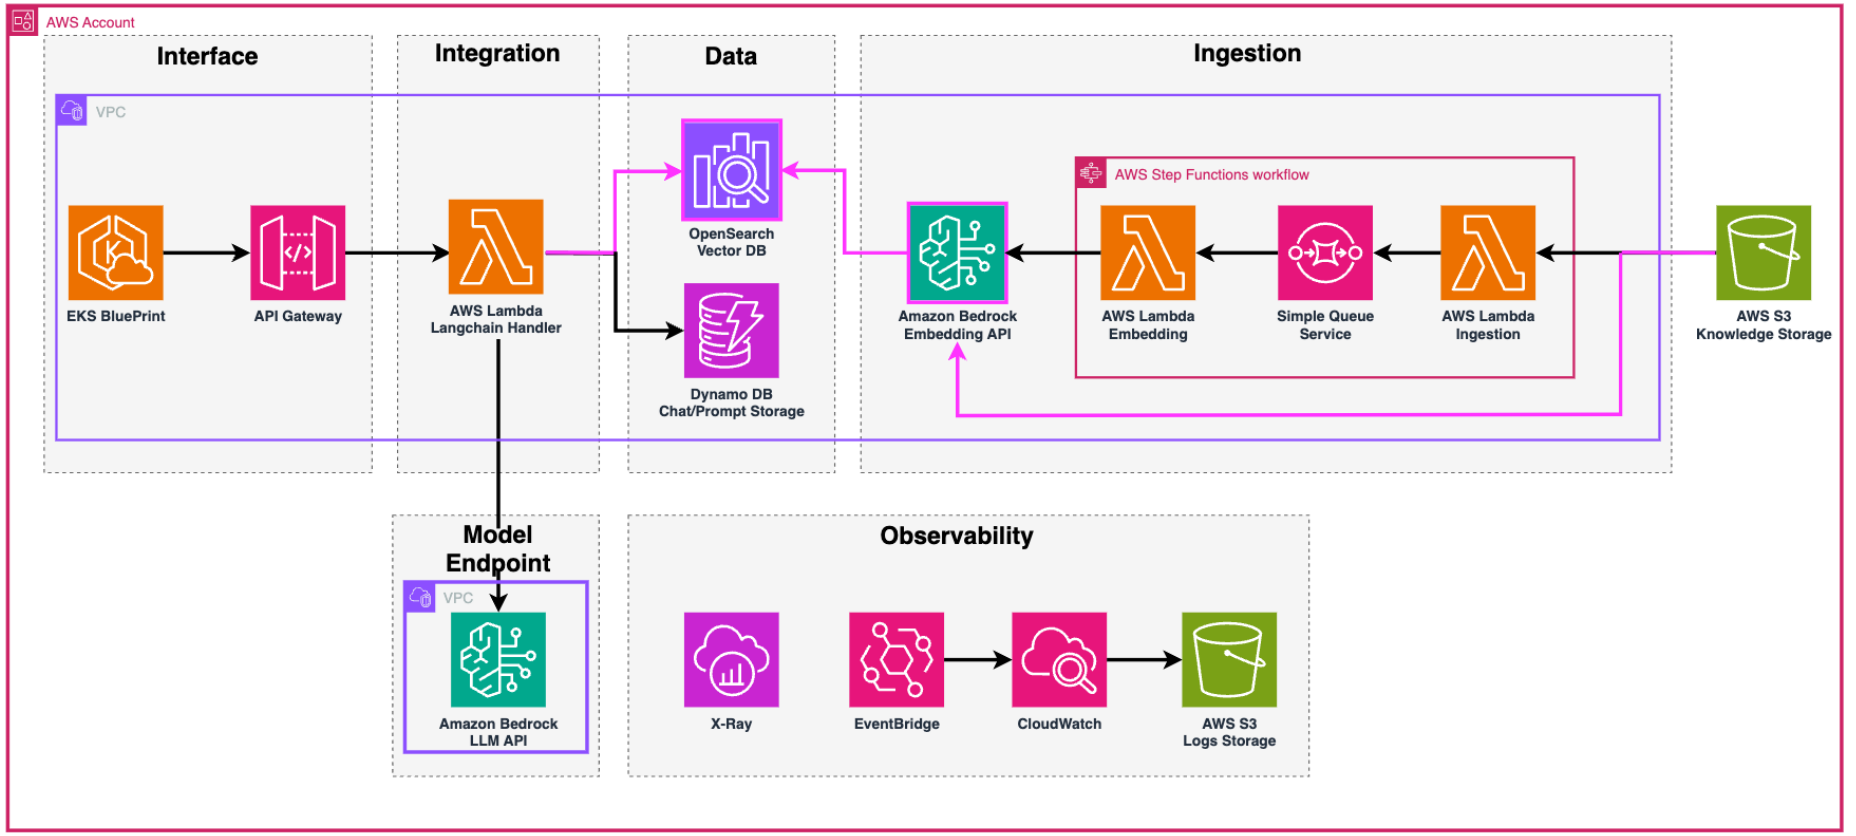

aws financial services

Alfred Discovery card

1. Use case objective:
   - Alfred is a GenAI application that provides intelligent document assistance for users to ask questions and receive answers based on ingested documents.
   - The success criteria are not defined yet, performance metrics is being defined
2. Data Ingestion:
   - Data types: Currently supports PDF files.
   - Data volume: Maximum 20 MB per document, no overall limit on the number of documents.
   - Data ingestion process: Manual import of documents.
   - Data organization: Concept of "folders" to group multiple documents as one.
3. Technical stack:
   - Front-end: React SPA with custom components for displaying answers.
   - Data storage: Local file storage for documents, Elasticsearch for indexing.
   - Language model: GPT-3.5 Turbo (chosen for speed over GPT-4 due to cost).
4. Challenges & evolutions:
   - OCR capability.
   - Support for C3, Word & PPT source data.
   - Table & Image handling.
   - Access to other source of data (internet, paid press like Bloomberg).
   - Performance: Improve upload performance (10-15 seconds for PDF text extraction)

Proposed ALFRED Architecture on AWS

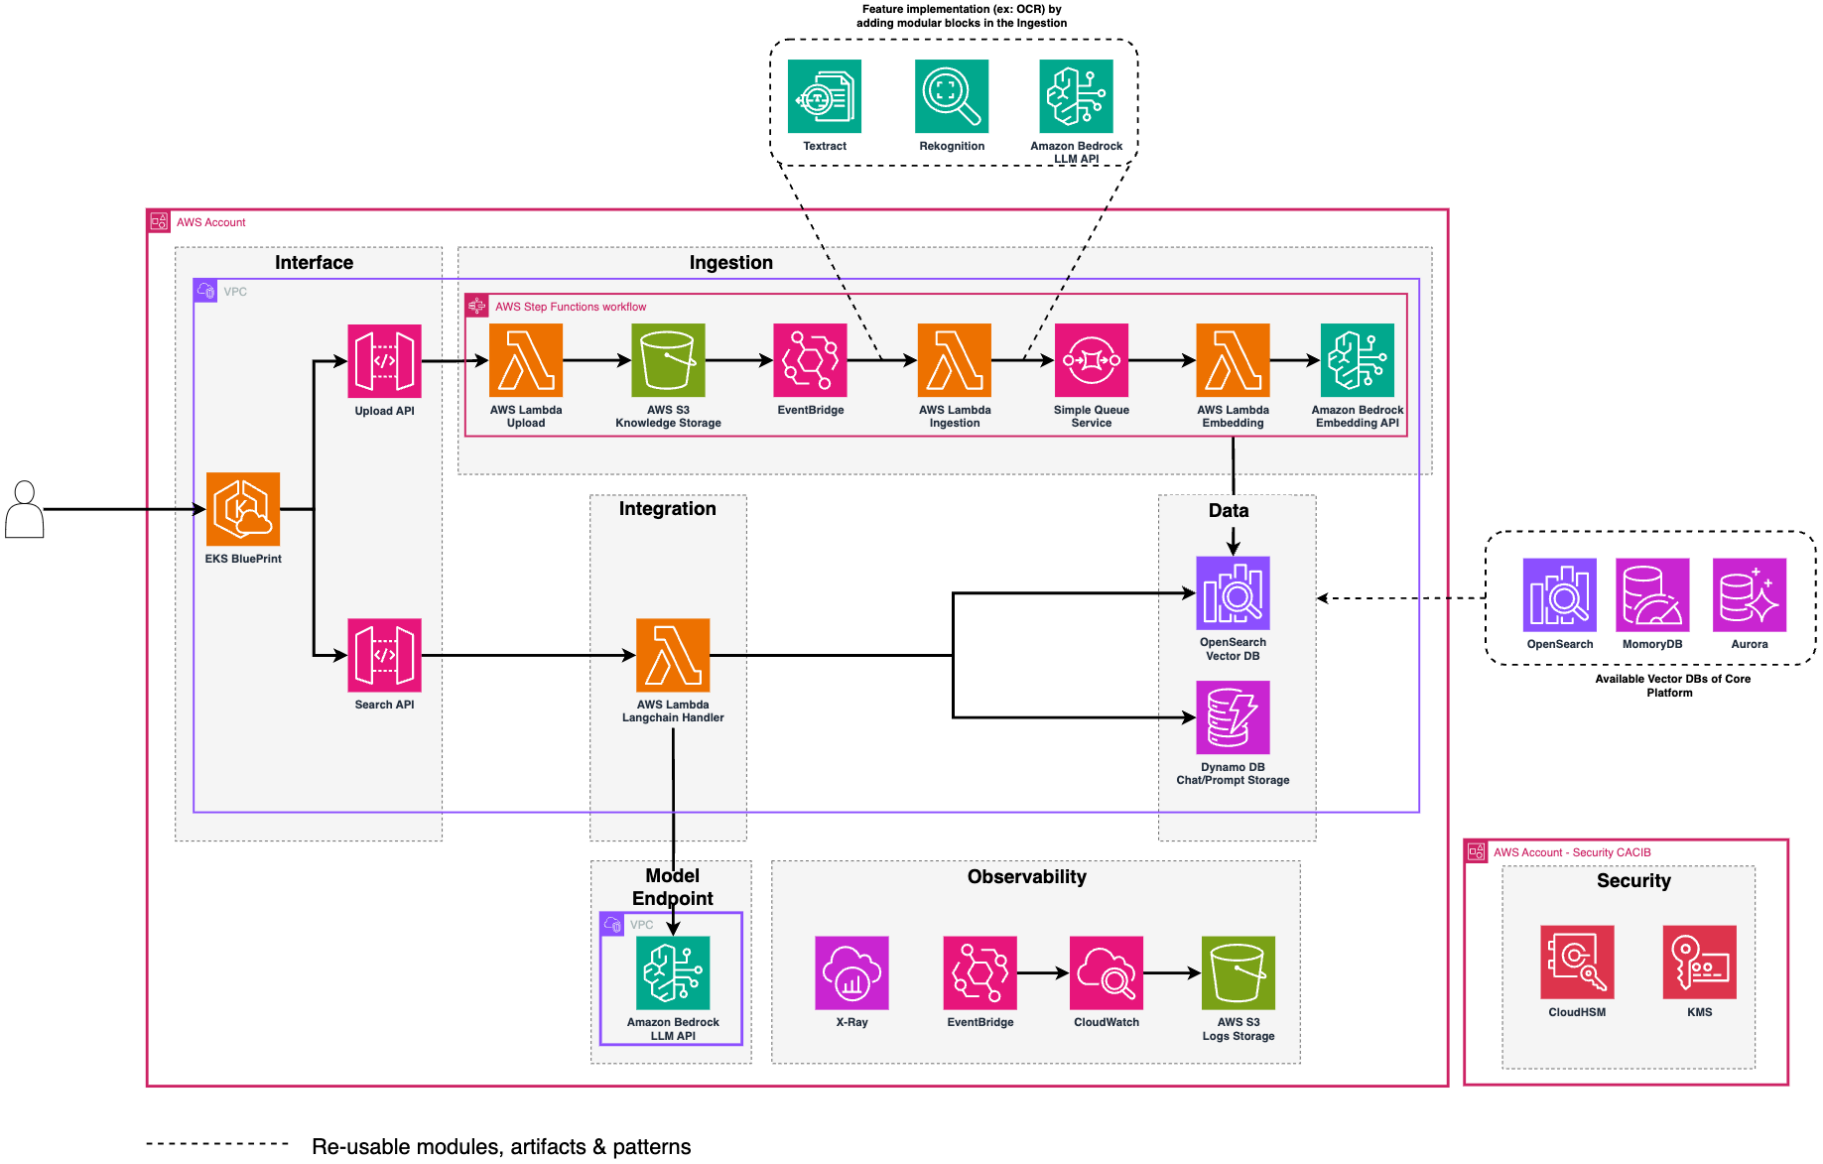

aws financial services

Augmented Sales Discovery card

1. Use case objective:
   - Augmented Sales support: Enhance the response time for sales teams to answer customer questions.
   - Success criteria: Less than 3 seconds response time with 300 users (front office).
2. Data Ingestion:
   - Data types: PDFs, with the ability to choose the source (only in English).
   - Data volume: 15 documents per day (max 20 pages), 1500 documents in batch.
   - Data sources: Bloomberg, Wiki (Ayca).
   - Ingestion frequency: Daily for Bloomberg, more frequent for Wiki.
3. Technical stack:
   - Selected model: Mistral 8x7b (based on chunk size and date range).
   - Index: Elastic Search.
   - UI: Streamlit.
4. Challenges & evolutions:
   - Data sources: Multiple data sources and segmentation based on user sign-on.
   - User feedback: Implement thumbs up/down and written feedback.
   - Document handling: Document splitter for tables and images (requirements, document analyzer developed internally).
   - Language support: Currently English, but planning to support French.
   - New features: Recommendation on past queries, suggestions.
   - Evaluation: Quantify the impact of hyperparameters (chunk size) on response quality and time.

aws financial services

Augmented Sales Architecture on AWS

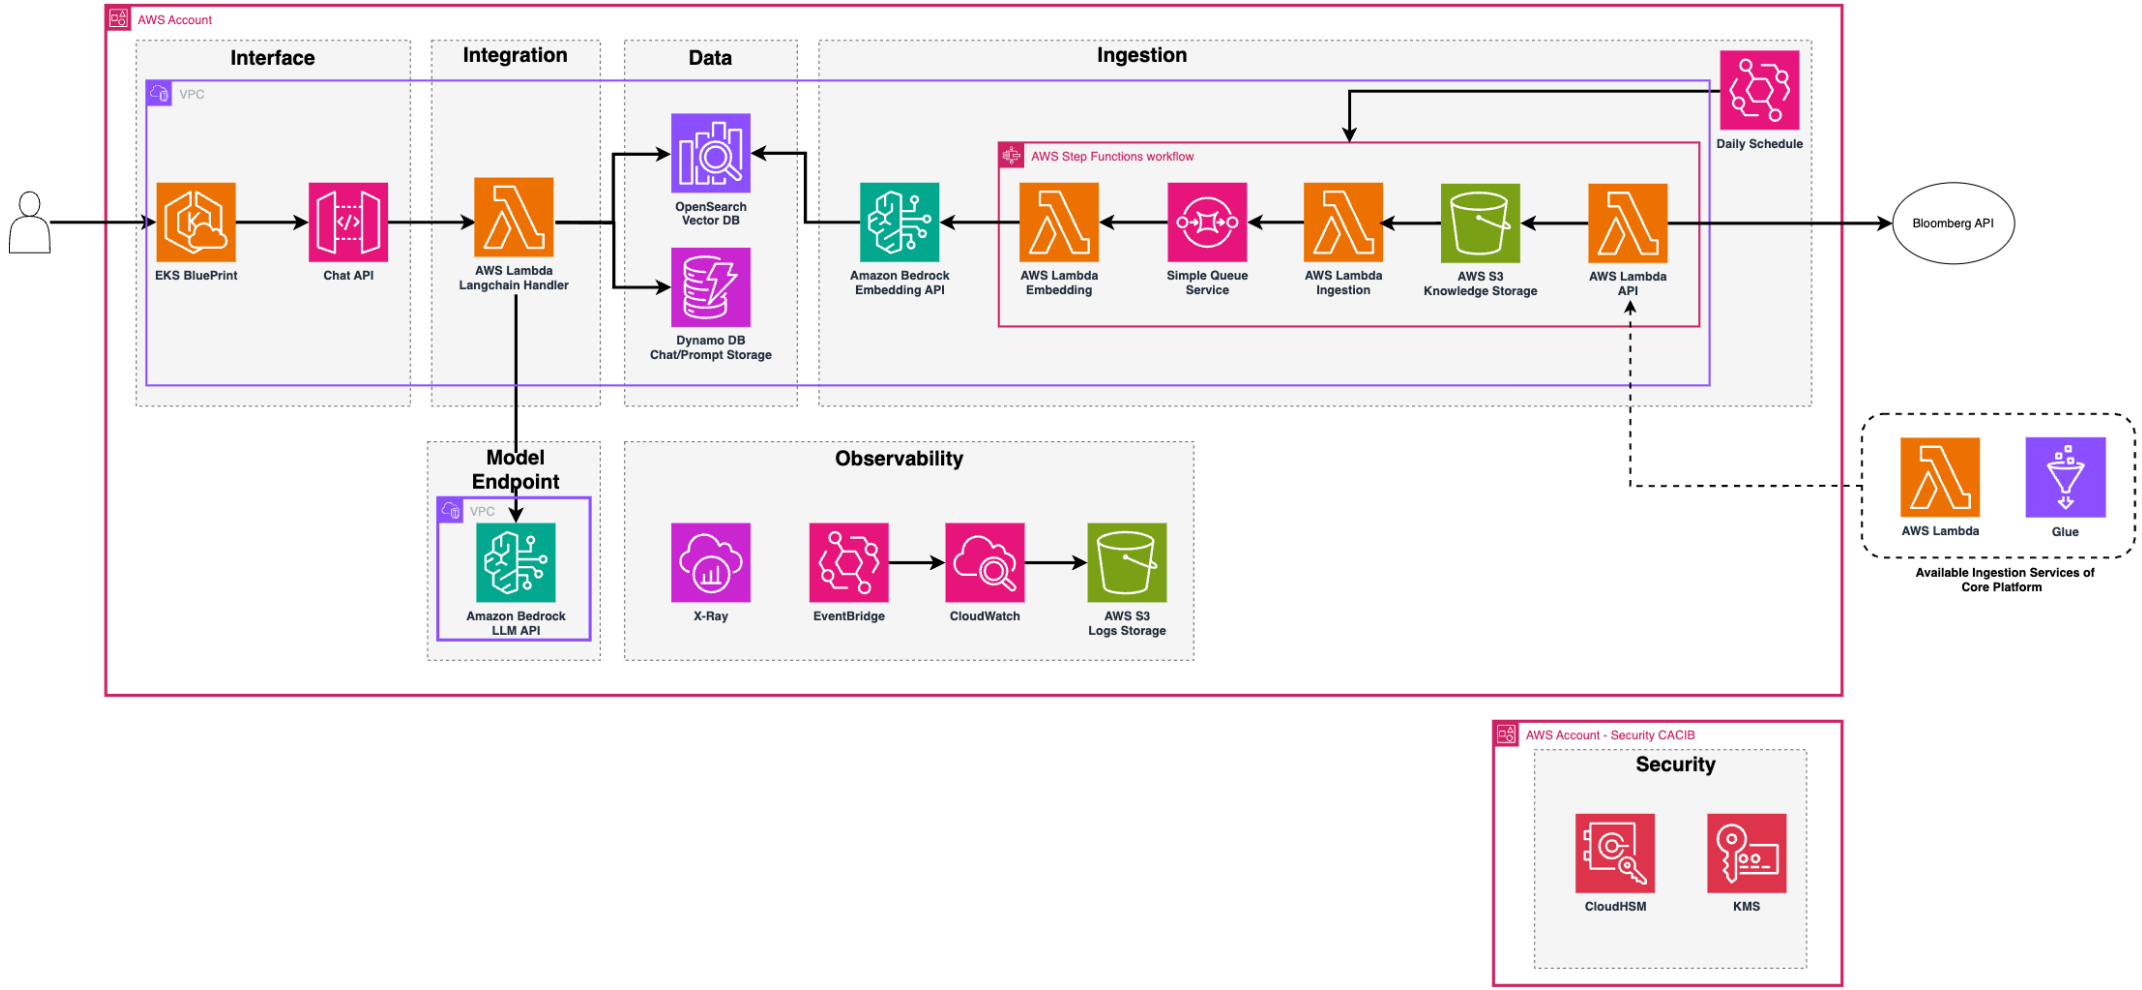

aws financial services

L1/L2 Support discovery card

1. Use case objective:
   - Technical Answer questions based on the target knowledge base
   - Generate pre-filled tickets
   - Provide a centralized access point for end-users to interact with a chatbot before reaching L1/L2 support
2. Data Ingestion:
   - 4 knowledge base sources, mainly from Confluence
   - Planned weekly batch ingestion
3. Technical stack:
   - Using LLaMA 2 and Naive RAG approach, no OpenAI integration
   - Using LangChain
   - FAISS vector DB et SQLite DB
4. Challenges and evolutions:
   - Table and image extraction, French documentation
   - Need to create labeled data to validate the use case
   - Wish for feedback loop, SSO authentication, UI, and performance monitoring

aws financial services

L1/L2 Support Architecture on AWS

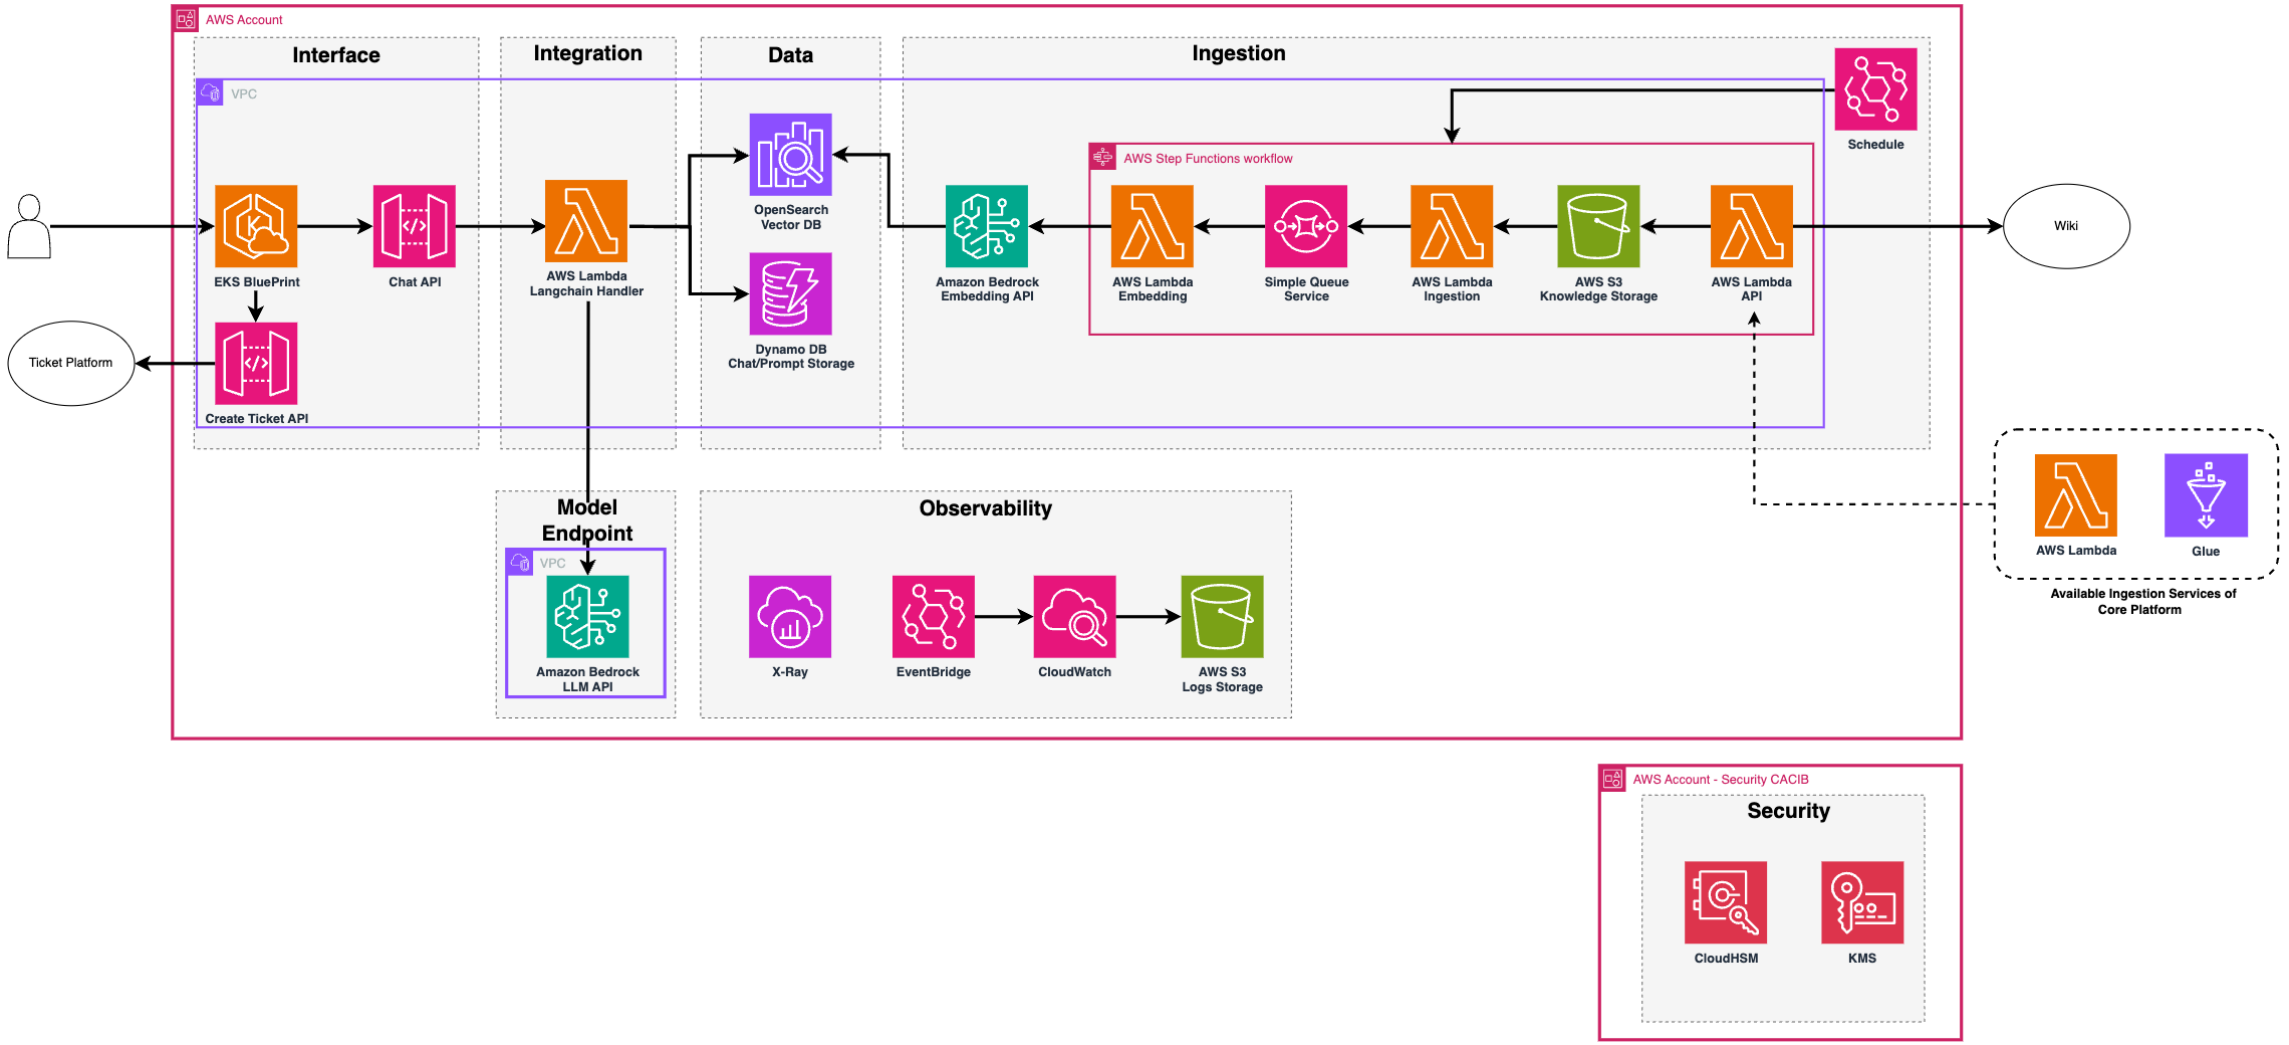

aws financial services

Market Colors Discovery discovery card

1. Use case objective:
   - The objective is to generate a daily newsletter/report for traders, combining public articles (news selection) and traders' insights.
   - The success criteria are to generate the report within approximately 30 minutes to 1 hour.
2. Data Ingestion:
   - Data sources: Structured data (Repo Data (KDB) and Market Data, Unstructured data (Web scraping from specific sites and user copy/paste)
   - Data ingestion frequency: Daily basis
3. Technical stack:
   - Data preparation Ingestion and ingestion: Dataiku (Data Science Platform - DSS), GPT
4. Challenges & evolutions:
   - Model insight generation improvement
   - Architecture decision - Hybrid architecture approach Front / Back remains On Premise leveraging Dataiku + Integration to AWS to access to LLM capabilities leveraging Amazon BedRock

Market Colors Architecture on AWS

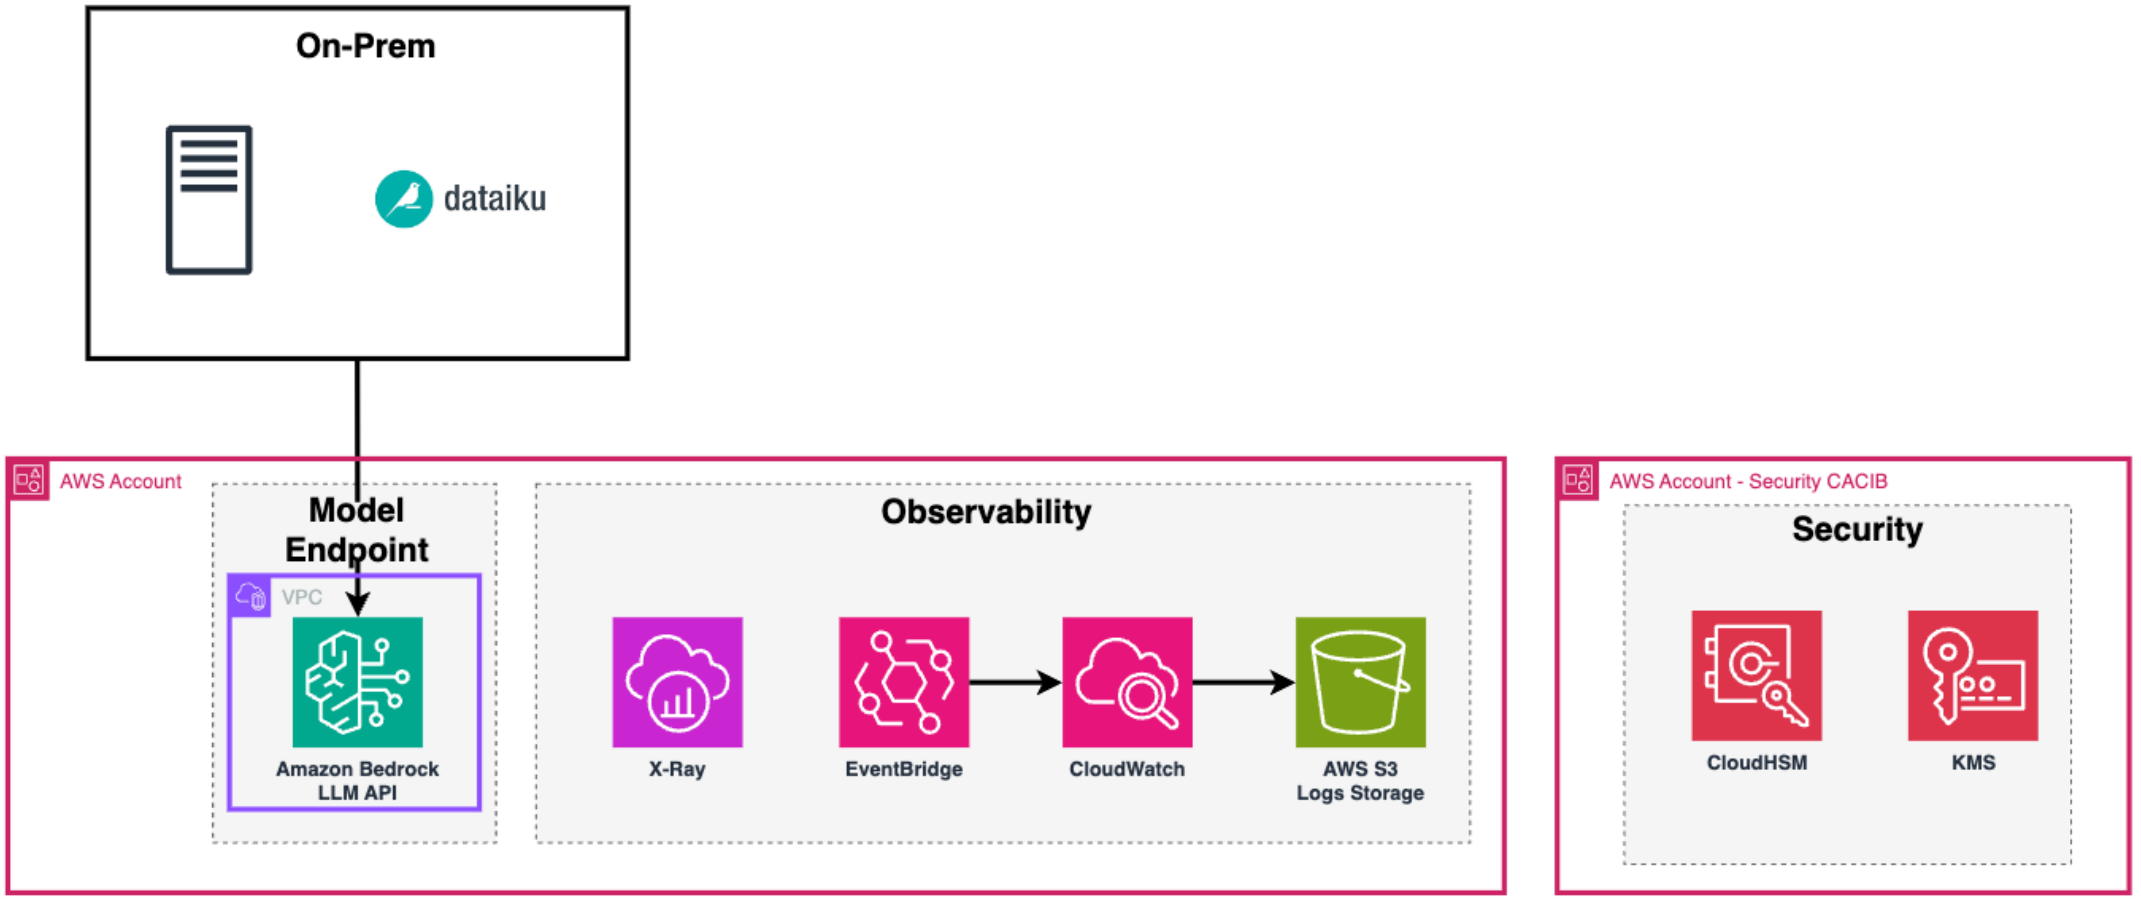

aws financial services

Appendix

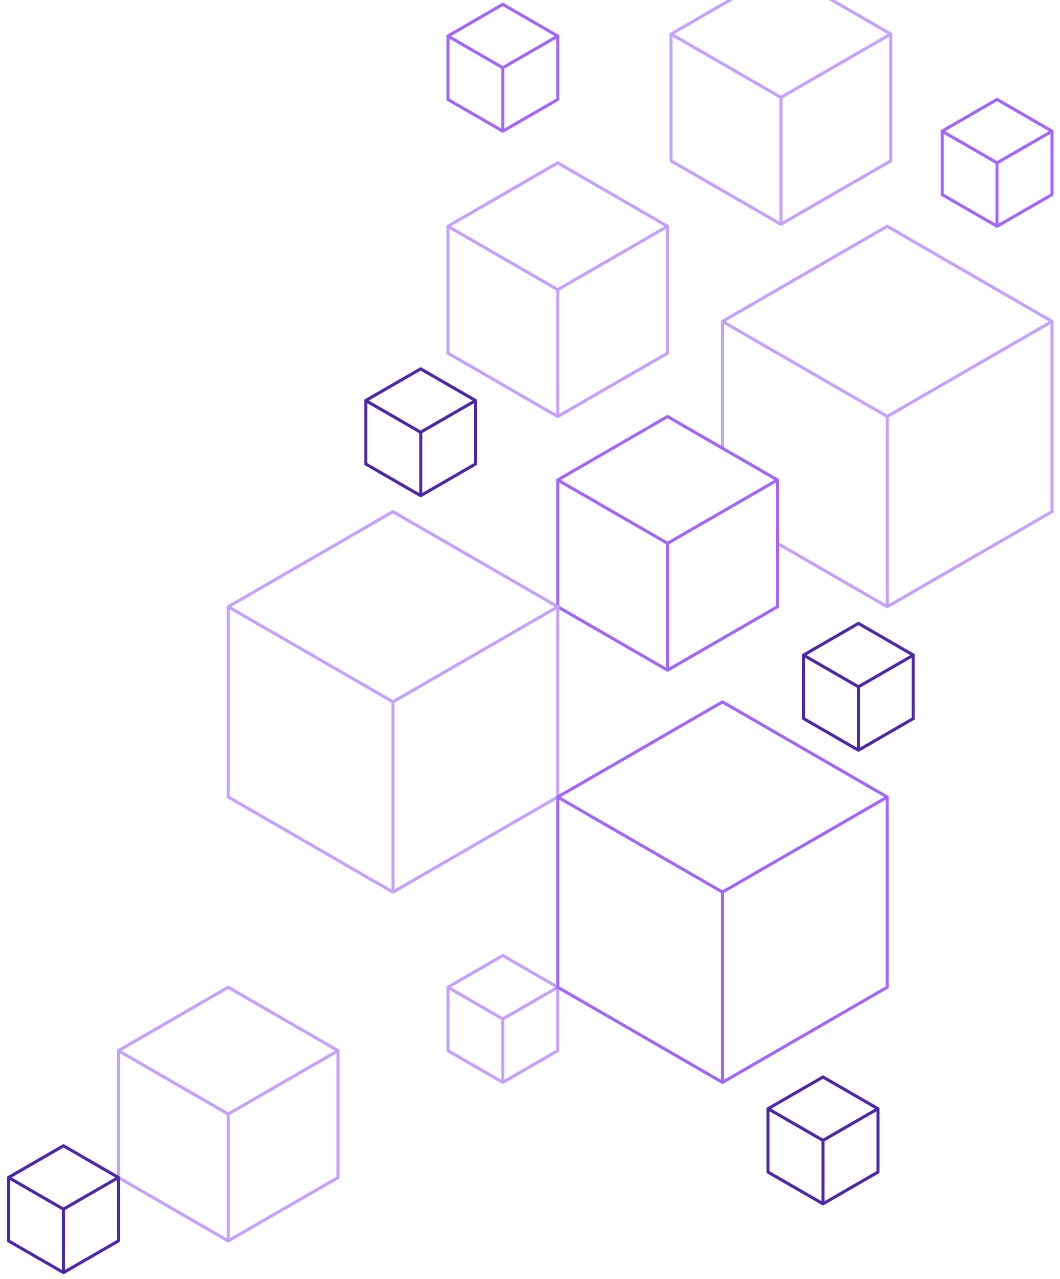

Centralized and Decentralized platform approaches

Decentralized platform

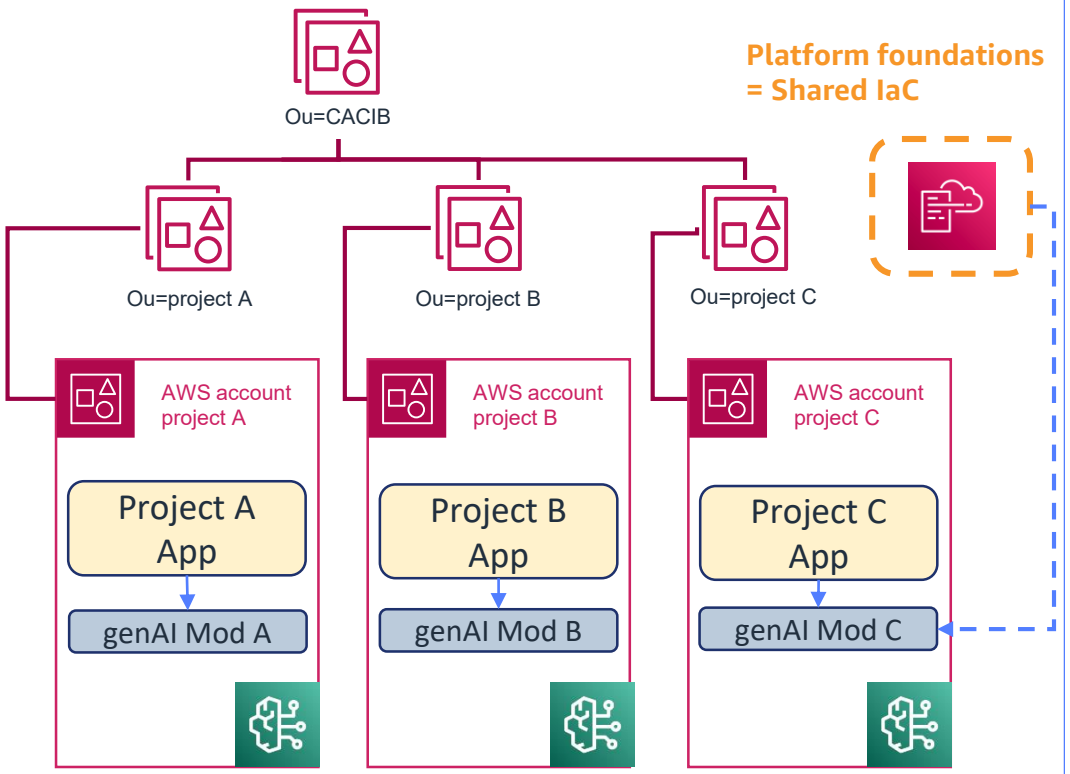

Centralized platform

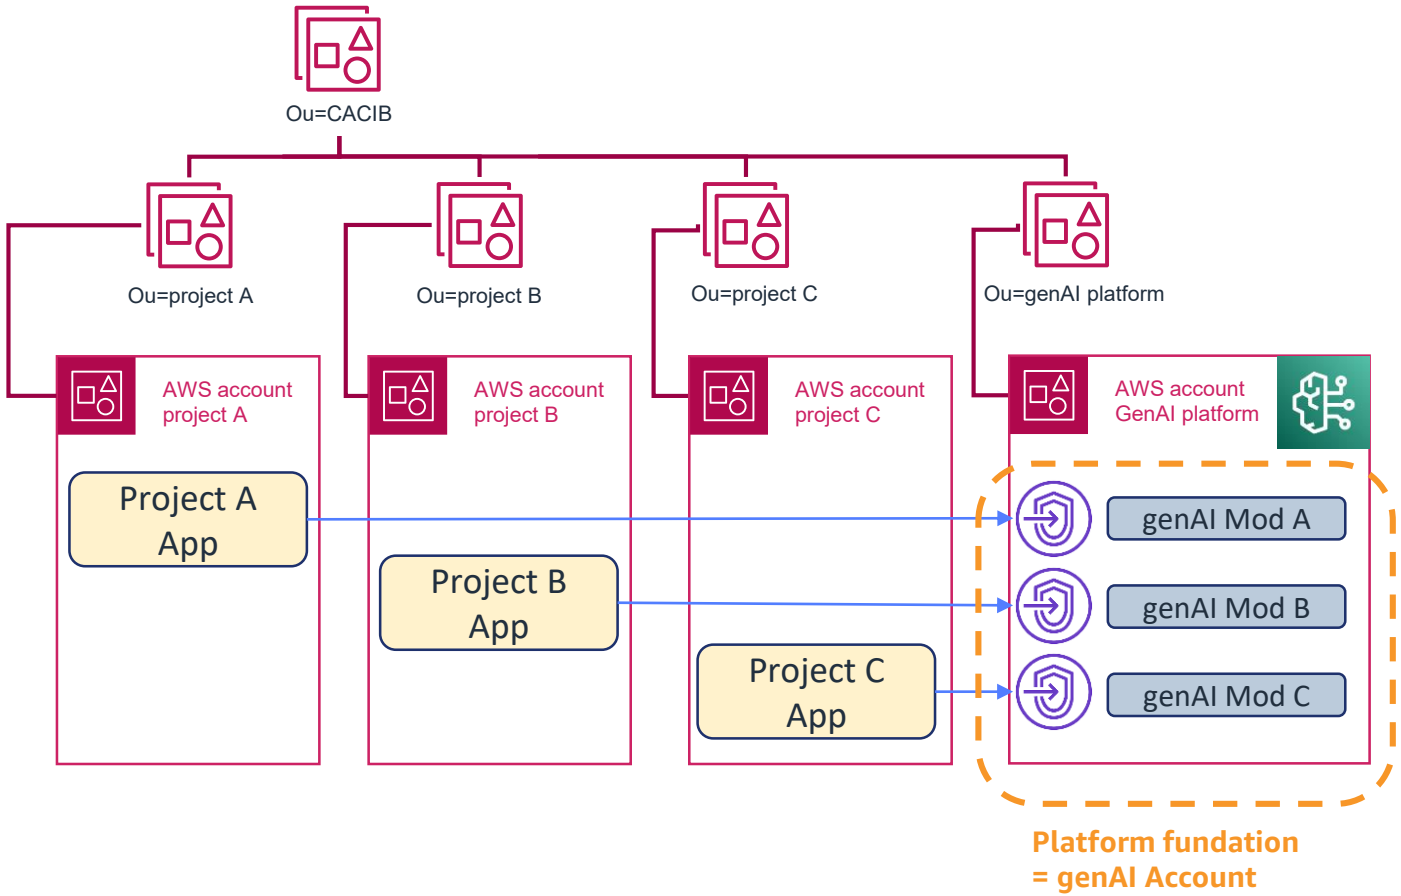

aws financial services

Centralized and Decentralized platform approaches

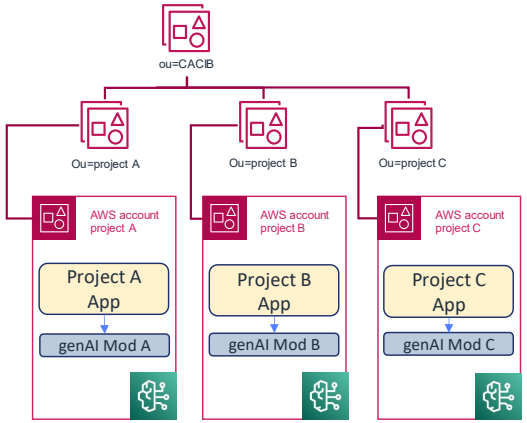

- Single tenant/project deployment
- Platform provides instantiated running services per project
- Platform centralized operations. Requires genAI platform cross project operations team.
- Architecture guidelines are enforced by DevOps practices and LZ policies
- genAI consumption structured per project according to LZ guidelines for general projects. FinOps is at the project level.
- Capabilities can be expanded for each project.

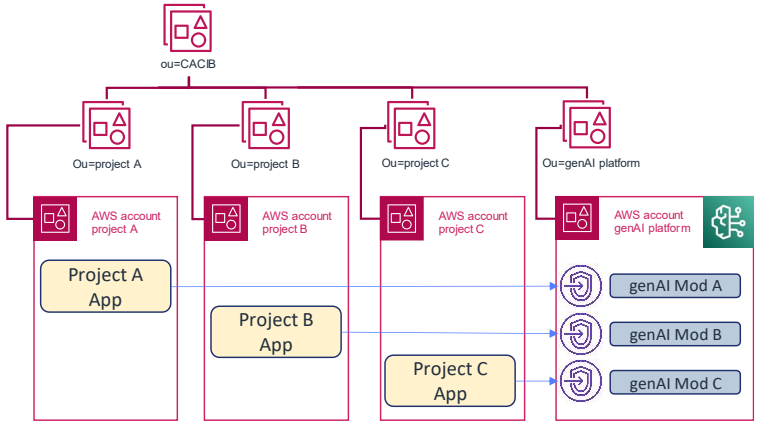

- Multi tenant/project deployment
- Platform provides shared running services
- Platform decentralized operations. Requires each project operations to integrate genAI component operations.
- Architecture guidelines are enforced by genAI platform project and DevOps practices.
- genAI consumption consolidated in a single account. Opportunity for cross projects finOps.
- Capabilities constrained by platform.

aws financial services

What are vector embeddings?

Source domain-specific data → **Tokenization** → **Vectorization** → **Store in vector data store** → **Perform semantic similarity search** → **Include semantically similar context in prompt**

Documents

Audio/video

Images

Semantic elements:
- Words, phrases
- Paragraphs, documents
- Scenes, song sections
- Faces, detected picture elements
- And more

0.35 0.1 0 0.9 001.0 00 0001.0 0 0...

0.35 0.1 0 0.8 001.0 00 0001.0 0 0...

0.15 0.1 0 0.7 001.0 00 0001.0 0 0...

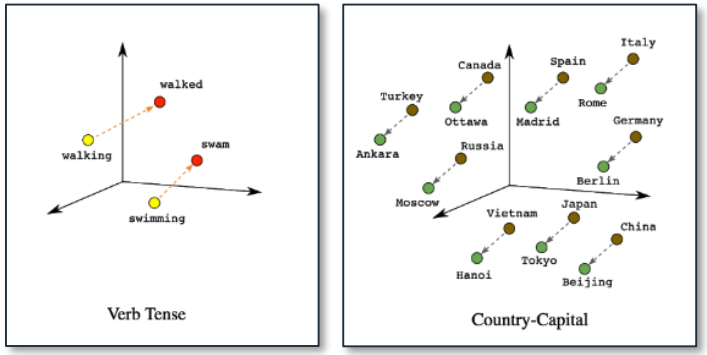

3D simplified representation. Embeddings can have thousands of dimensions. Source: https://daleonai.com/embeddings-explained

Embeddings: When vector elements are semantic, used in generative AI

aws financial services

A Retrieval-Augmented Generation (RAG) pattern

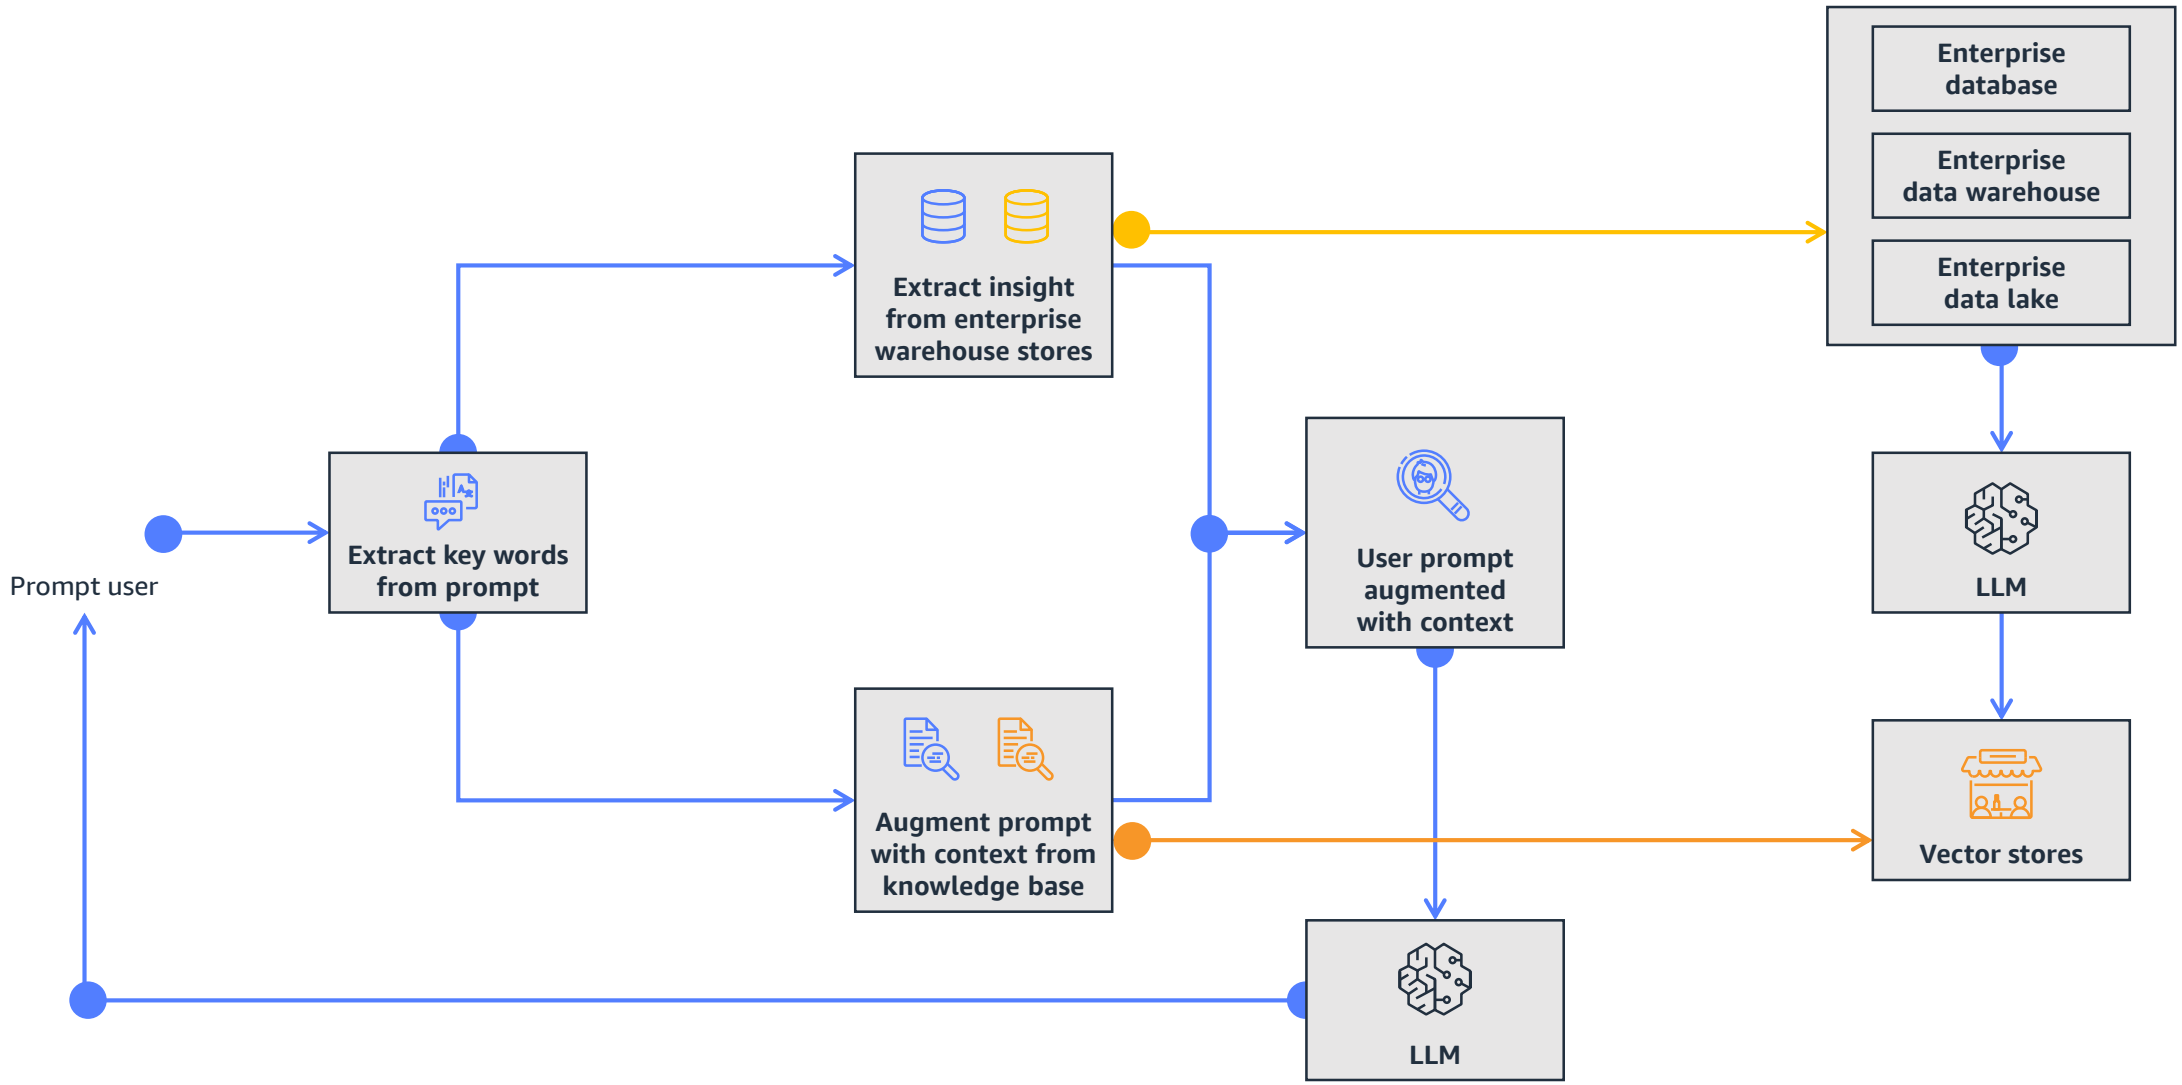

aws financial services

RAG reference architecture

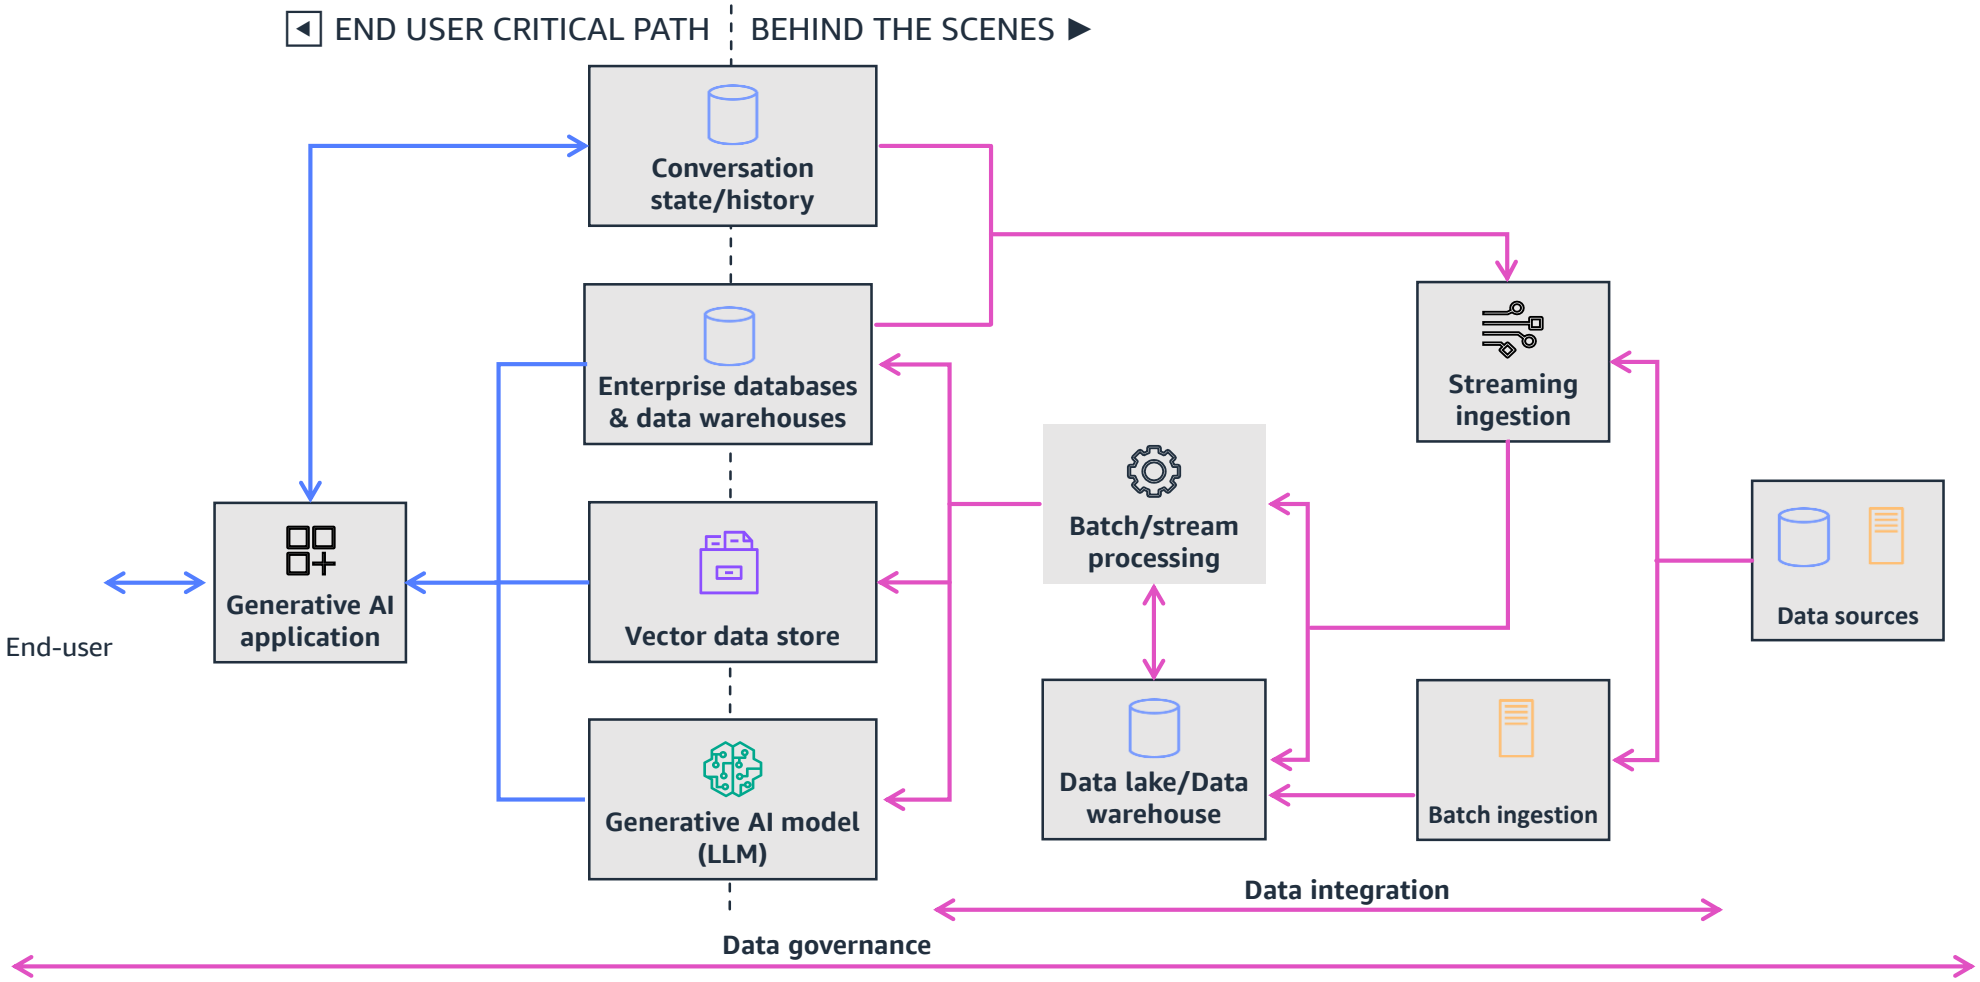

aws financial services

Data in the end user critical path

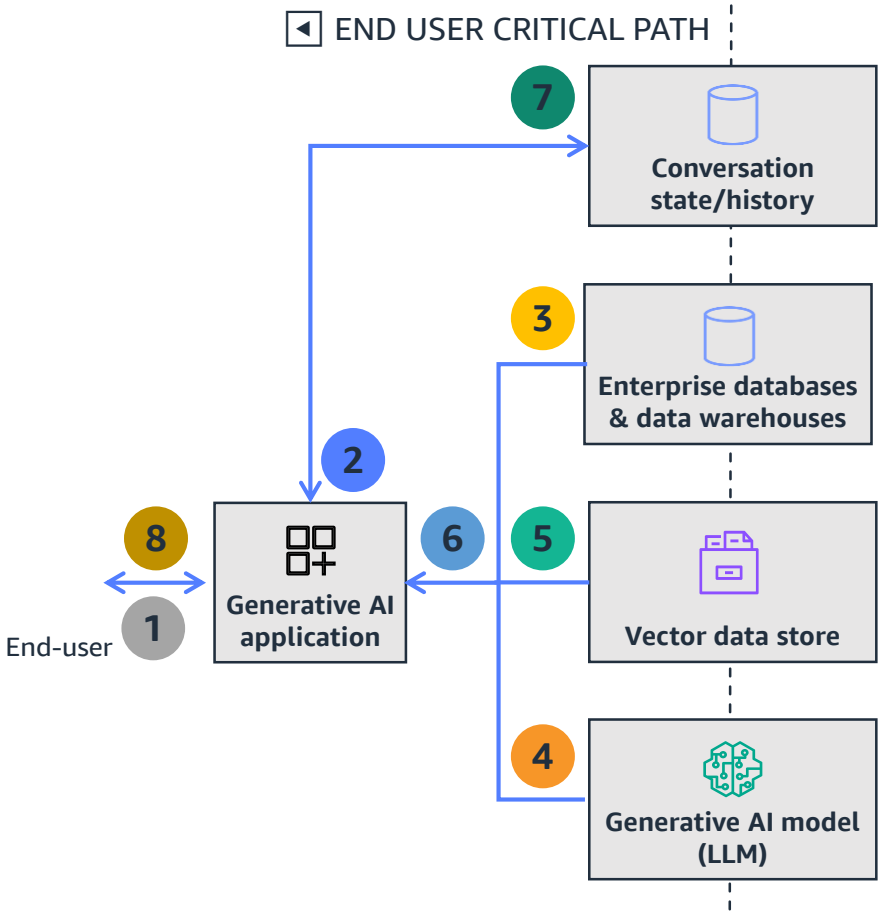

1. User asks question
2. Get conversation state/history
3. Get situational context
4. Tokenize, get question embedding from LLM
5. Run similarity search for question embedding, get top matching text
6. Invoke LLM w/ engineered prompt
7. Update conversation state/history
8. Return response

aws financial services

Knowledge bases for Amazon Bedrock

Fully managed native support for Retrieval Augmented Generation (RAG)

Automatically converts text documents into embeddings

Stores embeddings in your vector database

Retrieves embeddings and augments prompts

aws financial services

Vector databases and embedding models for Amazon Bedrock

Vector engine for Amazon OpenSearch Serverless

Redis Enterprise Cloud

Pinecone

Amazon Aurora

MongoDB

Amazon Titan Embeddings Text

Amazon Titan Multimodal Embeddings

Cohere Embed

aws financial services

Amazon Bedrock managed vs non-managed RAG

Amazon Bedrock **managed RAG**

- Automatic embeddings
- Managed embeddings in vector DB
- Automatic retrieval

- Limited embedding models
- Limited region availability
- A few control on retrieval prompts

Amazon Bedrock **non-managed RAG**

- Vast choice of embeddings models
- Several options of vector DB
- Custom retrieval prompts

- You might need to tune the embedding model
- You maintain the vector DB
- You integrate the systems

aws financial services